

Credit_Card_Fraud_Detection

December 12, 2024

Importing the modules

1 Welcome to the notebook

1.0.1 Task 1 - Importing the Dataset

```
[1]: import numpy as np
import pandas as pd

# Importing neural network modules
import tensorflow as tf
from tensorflow.keras.layers import Input, Dense, BatchNormalization, \
    LeakyReLU, Dropout
from tensorflow.keras.models import Model, Sequential
from tensorflow.keras.optimizers import Adam
from tensorflow.keras.initializers import RandomNormal

# Importing some machine learning modules
from sklearn.utils import shuffle
from sklearn.decomposition import PCA
from sklearn.preprocessing import StandardScaler

# Import data visualization modules
import seaborn as sns
import matplotlib.pyplot as plt
import plotly.express as px

print("Modules are imported!")
```

```
2024-12-12 12:20:14.010140: I tensorflow/core/platform/cpu_feature_guard.cc:193]
This TensorFlow binary is optimized with oneAPI Deep Neural Network Library
(oneDNN) to use the following CPU instructions in performance-critical
operations: SSE4.1 SSE4.2 AVX AVX2 AVX512F FMA
To enable them in other operations, rebuild TensorFlow with the appropriate
compiler flags.
```

Modules are imported!

```
[2]: data = pd.read_csv("Creditcard_dataset.csv")
data.head()
```

```
[2]:
```

	Time	V1	V2	V3	V4	V5	V6	V7	\
0	1	-0.966272	-0.185226	1.792993	-0.863291	-0.010309	1.247203	0.237609	
1	7	-0.644269	1.417964	1.074380	-0.492199	0.948934	0.428118	1.120631	
2	10	1.449044	-1.176339	0.913860	-1.375667	-1.971383	-0.629152	-1.423236	
3	10	0.384978	0.616109	-0.874300	-0.094019	2.924584	3.317027	0.470455	
4	11	1.069374	0.287722	0.828613	2.712520	-0.178398	0.337544	-0.096717	

	V8	V9	...	V21	V22	V23	V24	V25	\
0	0.377436	-1.387024	...	-0.108300	0.005274	-0.190321	-1.175575	0.647376	
1	-3.807864	0.615375	...	1.943465	-1.015455	0.057504	-0.649709	-0.415267	
2	0.048456	-1.720408	...	-0.009302	0.313894	0.027740	0.500512	0.251367	
3	0.538247	-0.558895	...	0.049924	0.238422	0.009130	0.996710	-0.767315	
4	0.115982	-0.221083	...	-0.036876	0.074412	-0.071407	0.104744	0.548265	

	V26	V27	V28	Amount	Class
0	-0.221929	0.062723	0.061458	123.50	0
1	-0.051634	-1.206921	-1.085339	40.80	0
2	-0.129478	0.042850	0.016253	7.80	0
3	-0.492208	0.042472	-0.054337	9.99	0
4	0.104094	0.021491	0.021293	27.50	0

[5 rows x 31 columns]

Check the data shape

```
[3]: data.shape
```

```
[3]: (50492, 31)
```

Let's see how many genuine and limited fraudulent records we have

```
[4]: data.Class.value_counts()
```

```
[4]: 0    50000
      1     492
      Name: Class, dtype: int64
```

1.0.2 Task 2 - Data Preprocessing and Exploration

- Removing all the rows with Nan values
- Removing Time column
- Feature Scaling Amount column
- Split the data into features and labels
- Data Exploration

Removing the rows Nan values in the dataset

```
[5]: data.dropna(inplace=True)
data.shape
```

```
[5]: (50492, 31)
```

Removing Time column

```
[8]: data = data.drop(axis = 1, columns = 'Time')
data.head()
```

```
[8]:
```

	V1	V2	V3	V4	V5	V6	V7 \
0	-0.966272	-0.185226	1.792993	-0.863291	-0.010309	1.247203	0.237609
1	-0.644269	1.417964	1.074380	-0.492199	0.948934	0.428118	1.120631
2	1.449044	-1.176339	0.913860	-1.375667	-1.971383	-0.629152	-1.423236
3	0.384978	0.616109	-0.874300	-0.094019	2.924584	3.317027	0.470455
4	1.069374	0.287722	0.828613	2.712520	-0.178398	0.337544	-0.096717

	V8	V9	V10	...	V21	V22	V23	V24 \
0	0.377436	-1.387024	-0.054952	...	-0.108300	0.005274	-0.190321	-1.175575
1	-3.807864	0.615375	1.249376	...	1.943465	-1.015455	0.057504	-0.649709
2	0.048456	-1.720408	1.626659	...	-0.009302	0.313894	0.027740	0.500512
3	0.538247	-0.558895	0.309755	...	0.049924	0.238422	0.009130	0.996710
4	0.115982	-0.221083	0.460230	...	-0.036876	0.074412	-0.071407	0.104744

	V25	V26	V27	V28	Amount	Class
0	0.647376	-0.221929	0.062723	0.061458	123.50	0
1	-0.415267	-0.051634	-1.206921	-1.085339	40.80	0
2	0.251367	-0.129478	0.042850	0.016253	7.80	0
3	-0.767315	-0.492208	0.042472	-0.054337	9.99	0
4	0.548265	0.104094	0.021491	0.021293	27.50	0

[5 rows x 30 columns]

Feature Scaling of Amount column

```
[13]: scaler = StandardScaler()
data['Amount'] = scaler.fit_transform(data[['Amount']])
data.head()
```

```
[13]:
```

	V1	V2	V3	V4	V5	V6	V7 \
0	-0.966272	-0.185226	1.792993	-0.863291	-0.010309	1.247203	0.237609
1	-0.644269	1.417964	1.074380	-0.492199	0.948934	0.428118	1.120631
2	1.449044	-1.176339	0.913860	-1.375667	-1.971383	-0.629152	-1.423236
3	0.384978	0.616109	-0.874300	-0.094019	2.924584	3.317027	0.470455
4	1.069374	0.287722	0.828613	2.712520	-0.178398	0.337544	-0.096717

	V8	V9	V10	...	V21	V22	V23	V24 \
0	0.377436	-1.387024	-0.054952	...	-0.108300	0.005274	-0.190321	-1.175575

1	-3.807864	0.615375	1.249376	...	1.943465	-1.015455	0.057504	-0.649709
2	0.048456	-1.720408	1.626659	...	-0.009302	0.313894	0.027740	0.500512
3	0.538247	-0.558895	0.309755	...	0.049924	0.238422	0.009130	0.996710
4	0.115982	-0.221083	0.460230	...	-0.036876	0.074412	-0.071407	0.104744

	V25	V26	V27	V28	Amount	Class
0	0.647376	-0.221929	0.062723	0.061458	0.150105	0
1	-0.415267	-0.051634	-1.206921	-1.085339	-0.199848	0
2	0.251367	-0.129478	0.042850	0.016253	-0.339490	0
3	-0.767315	-0.492208	0.042472	-0.054337	-0.330223	0
4	0.548265	0.104094	0.021491	0.021293	-0.256128	0

[5 rows x 30 columns]

Let's split the genuine and fraud records into separate dataframes

```
[14]: data_fraud = data[data.Class == 1]
      data_genuine = data[data.Class == 0]

      data_fraud
```

```
[14]:
```

	V1	V2	V3	V4	V5	V6	V7 \
50000	-2.312227	1.951992	-1.609851	3.997906	-0.522188	-1.426545	-2.537387
50001	-3.043541	-3.157307	1.088463	2.288644	1.359805	-1.064823	0.325574
50002	-2.303350	1.759247	-0.359745	2.330243	-0.821628	-0.075788	0.562320
50003	-4.397974	1.358367	-2.592844	2.679787	-1.128131	-1.706536	-3.496197
50004	1.234235	3.019740	-4.304597	4.732795	3.624201	-1.357746	1.713445
...
50487	-1.927883	1.125653	-4.518331	1.749293	-1.566487	-2.010494	-0.882850
50488	1.378559	1.289381	-5.004247	1.411850	0.442581	-1.326536	-1.413170
50489	-0.676143	1.126366	-2.213700	0.468308	-1.120541	-0.003346	-2.234739
50490	-3.113832	0.585864	-5.399730	1.817092	-0.840618	-2.943548	-2.208002
50491	1.991976	0.158476	-2.583441	0.408670	1.151147	-0.096695	0.223050

	V8	V9	V10	...	V21	V22	V23 \
50000	1.391657	-2.770089	-2.772272	...	0.517232	-0.035049	-0.465211
50001	-0.067794	-0.270953	-0.838587	...	0.661696	0.435477	1.375966
50002	-0.399147	-0.238253	-1.525412	...	-0.294166	-0.932391	0.172726
50003	-0.248778	-0.247768	-4.801637	...	0.573574	0.176968	-0.436207
50004	-0.496358	-1.282858	-2.447469	...	-0.379068	-0.704181	-0.656805
...
50487	0.697211	-2.064945	-5.587794	...	0.778584	-0.319189	0.639419
50488	0.248525	-1.127396	-3.232153	...	0.370612	0.028234	-0.145640
50489	1.210158	-0.652250	-3.463891	...	0.751826	0.834108	0.190944
50490	1.058733	-1.632333	-5.245984	...	0.583276	-0.269209	-0.456108
50491	-0.068384	0.577829	-0.888722	...	-0.164350	-0.295135	-0.072173

	V24	V25	V26	V27	V28	Amount	Class
50000	0.320198	0.044519	0.177840	0.261145	-0.143276	-0.372497	1
50001	-0.293803	0.279798	-0.145362	-0.252773	0.035764	1.866017	1
50002	-0.087330	-0.156114	-0.542628	0.039566	-0.153029	0.642790	1
50003	-0.053502	0.252405	-0.657488	-0.827136	0.849573	-0.122833	1
50004	-1.632653	1.488901	0.566797	-0.010016	0.146793	-0.368265	1
...
50487	-0.294885	0.537503	0.788395	0.292680	0.147968	1.277825	1
50488	-0.081049	0.521875	0.739467	0.389152	0.186637	-0.369281	1
50489	0.032070	-0.739695	0.471111	0.385107	0.194361	-0.042898	1
50490	-0.183659	-0.328168	0.606116	0.884876	-0.253700	0.664244	1
50491	-0.450261	0.313267	-0.289617	0.002988	-0.015309	-0.192527	1

[492 rows x 30 columns]

Split the data into features and labels

```
[17]: X = data.drop("Class", axis = 1)
      y = data.Class
```

Data Exploration - Apply PCA to reduce the dimensionality of features X into two dimensions -
Use a scatter plot to visualize our data

```
[25]: pca = PCA(2)
      transformed_data = pca.fit_transform(X)
      df = pd.DataFrame(transformed_data)
      df['label'] = y
      df
```

```
[25]:
```

	0	1	label
0	-0.447839	1.197491	0
1	-0.582394	0.258152	0
2	-0.939386	-0.728332	0
3	-0.630768	-0.499062	0
4	-0.536287	-1.055405	0
...
50487	10.855754	-2.787413	1
50488	6.927758	-4.949699	1
50489	7.038194	-2.785052	1
50490	11.649180	-1.991794	1
50491	-0.227982	-2.009630	1

[50492 rows x 3 columns]

Let's Use a scatter plot to visualize our data

```
[27]: px.scatter(df , x = 0, y = 1, color = df.label.astype(str))
```

1.0.3 Task 3 - Building the Generator Model

Write a method to create the Generator model architecture

```
[28]: def build_generator():
    model = Sequential()

    model.add(Dense(32, activation='relu', input_dim=29,
kernel_initializer='he_uniform'))
    model.add(BatchNormalization())

    model.add(Dense(64, activation='relu'))
    model.add(BatchNormalization())

    model.add(Dense(128, activation='relu'))
    model.add(BatchNormalization())

    model.add(Dense(29, activation='relu'))
    model.summary()

    return model
build_generator()
```

```
2024-12-12 12:59:58.593407: I tensorflow/core/platform/cpu_feature_guard.cc:193]
This TensorFlow binary is optimized with oneAPI Deep Neural Network Library
(oneDNN) to use the following CPU instructions in performance-critical
operations: SSE4.1 SSE4.2 AVX AVX2 AVX512F FMA
To enable them in other operations, rebuild TensorFlow with the appropriate
compiler flags.
```

Model: "sequential"

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 32)	960
batch_normalization (BatchNormalization)	(None, 32)	128
dense_1 (Dense)	(None, 64)	2112
batch_normalization_1 (BatchNormalization)	(None, 64)	256
dense_2 (Dense)	(None, 128)	8320
batch_normalization_2 (BatchNormalization)	(None, 128)	512

dense_3 (Dense)	(None, 29)	3741
-----------------	------------	------

```
=====
Total params: 16,029
Trainable params: 15,581
Non-trainable params: 448
-----
```

[28]: <keras.engine.sequential.Sequential at 0x7c94412b7cd0>

1.0.4 Task 4 - Building the Discriminator Model

Write a method to create the Discriminator model architecture

```
[30]: def build_discriminator():
    model = Sequential()

    # Correcting the parameter name to match TensorFlow/Keras requirements
    model.add(Dense(128, activation='relu', input_dim=29,
↪kernel_initializer='he_uniform'))
    model.add(Dense(64, activation='relu'))
    model.add(Dense(32, activation='relu'))
    model.add(Dense(1, activation='sigmoid'))

    # Compile the model with appropriate loss and optimizer
    model.compile(optimizer='adam', loss='binary_crossentropy')

    return model

# Build the discriminator
discriminator = build_discriminator()
discriminator.summary()
```

Model: "sequential_2"

Layer (type)	Output Shape	Param #
dense_4 (Dense)	(None, 128)	3840
dense_5 (Dense)	(None, 64)	8256
dense_6 (Dense)	(None, 32)	2080
dense_7 (Dense)	(None, 1)	33

```
=====
Total params: 14,209
Trainable params: 14,209
```

Non-trainable params: 0

1.0.5 Task 5 - Combine Generator and Discriminator models to Build The GAN

```
[31]: def build_gan(generator, discriminator):
        discriminator.trainable = False
        gan_input = Input(shape = (generator.input_shape[1],))
        x = generator(gan_input)

        gan_output = discriminator(x)
        gan = Model(gan_input, gan_output)
        gan.summar()
        return gan
```

Let's create a method that generates synthetic data using the Generator

```
[32]: def generate_synthetic_data(generator, num_samples):
        noise = np.random.normal(0, 1, (num_samples, generator.input_shape[1]))
        fake_data = generator.predict(noise)
        return fake_data
```

1.0.6 Task 6 - Train and evaluate our GAN

- Defining some variables
- Creating our GAN
- Training the GAN
- Monitor the GAN performance using PCA

```
[33]: def monitor_generator(generator):
        # Initialize a PCA (Principal Component Analysis) object with 2 components
        pca = PCA(n_components=2)

        # Drop the 'Class' column from the fraud dataset to get real data
        real_fraud_data = data_fraud.drop("Class", axis=1)

        # Transform the real fraud data using PCA
        transformed_data_real = pca.fit_transform(real_fraud_data.values)

        # Create a DataFrame for the transformed real data and add a 'label' column
        ↪with the value 'real'
        df_real = pd.DataFrame(transformed_data_real)
        df_real['label'] = "real"

        # Generate synthetic fraud data using the provided generator and specify
        ↪the number of samples (492 in this case)
        synthetic_fraud_data = generate_synthetic_data(generator, 492)
```



```

# Transform the synthetic fraud data using PCA
transformed_data_fake = pca.fit_transform(synthetic_fraud_data)

# Create a DataFrame for the transformed fake data and add a 'label' column
↳with the value 'fake'
df_fake = pd.DataFrame(transformed_data_fake)
df_fake['label'] = "fake"

# Concatenate the real and fake data DataFrames
df_combined = pd.concat([df_real, df_fake])

# Create a scatterplot to visualize the data points, using the first and
↳second PCA components as x and y, respectively,
# and color points based on the 'label' column, with a size of 10
plt.figure()
sns.scatterplot(data=df_combined, x=0, y=1, hue='label', s=10)
plt.show()

```

```

[39]: from tensorflow.keras.models import Model
from tensorflow.keras.layers import Input

def build_gan(generator, discriminator):
    # Freeze the discriminator's weights during GAN training
    discriminator.trainable = False

    # Create input for the GAN
    gan_input = Input(shape=(generator.input_shape[1],))
    gan_output = discriminator(generator(gan_input)) # Pass through generator
    ↳and discriminator
    gan = Model(gan_input, gan_output)

    return gan

# Build and compile the GAN
generator = build_generator()
discriminator = build_discriminator()
gan = build_gan(generator, discriminator)
gan.compile(optimizer='adam', loss='binary_crossentropy')

# Display the GAN model architecture
gan.summary()

```

Model: "sequential_7"

Layer (type)	Output Shape	Param #
dense_24 (Dense)	(None, 32)	960

```

batch_normalization_9 (Batch Normalization) (None, 32) 128
dense_25 (Dense) (None, 64) 2112
batch_normalization_10 (Batch Normalization) (None, 64) 256
dense_26 (Dense) (None, 128) 8320
batch_normalization_11 (Batch Normalization) (None, 128) 512
dense_27 (Dense) (None, 29) 3741
=====
Total params: 16,029
Trainable params: 15,581
Non-trainable params: 448
-----
Model: "model_2"
-----
Layer (type)                Output Shape              Param #
-----
input_3 (InputLayer)        [(None, 29)]              0
sequential_7 (Sequential)   (None, 29)                16029
sequential_8 (Sequential)   (None, 1)                14209
=====
Total params: 30,238
Trainable params: 15,581
Non-trainable params: 14,657
-----

```

1.0.7 Task 7 - Generate synthetic data using the trained Generator

- Generate 1000 fraudulent data points using the trained generator
- Compare the distribution of **real** and **synthetic** fraudulent data points.

```

[40]: import pandas as pd
import numpy as np
import plotly.express as px
import matplotlib.pyplot as plt

# Generate synthetic data

```

```

def generate_and_compare(generator, real_data, num_samples=1000):
    # Generate synthetic data using the trained generator
    synthetic_data = generate_synthetic_data(generator, num_samples)

    # Create DataFrame for synthetic data and label it as 'fake'
    df_synthetic = pd.DataFrame(synthetic_data, columns=real_data.columns)
    df_synthetic['label'] = 'fake'

    # Label the real data as 'real'
    df_real = real_data.copy()
    df_real['label'] = 'real'

    # Combine real and synthetic data into one DataFrame
    combined_df = pd.concat([df_real, df_synthetic], ignore_index=True)

    # Return the combined DataFrame for further analysis
    return combined_df

# Generate combined data
real_data = data_fraud.drop("Class", axis=1) # Drop target column from real_
↳data
combined_df = generate_and_compare(generator, real_data)

```

32/32 [=====] - 0s 887us/step

Checking the individual feature distribution of synthetic and real fraud data.

```

[41]: for col in combined_df.columns:
    plt.figure()
    fig = px.histogram(combined_df, color = 'label', x=col, barmode="overlay",
    ↳title = f'Feature {col}', width = 640, height = 500)
    fig.show()

```

<Figure size 640x480 with 0 Axes>

<Figure size 640x480 with 0 Axes>

<Figure size 640x480 with 0 Axes>

<Figure size 640x480 with 0 Axes>

<Figure size 640x480 with 0 Axes>

<Figure size 640x480 with 0 Axes>

<Figure size 640x480 with 0 Axes>

<Figure size 640x480 with 0 Axes>

<Figure size 640x480 with 0 Axes>

<Figure size 640x480 with 0 Axes>

<Figure size 640x480 with 0 Axes>
<Figure size 640x480 with 0 Axes>
<Figure size 640x480 with 0 Axes>
<Figure size 640x480 with 0 Axes>
<Figure size 640x480 with 0 Axes>
<Figure size 640x480 with 0 Axes>
<Figure size 640x480 with 0 Axes>
<Figure size 640x480 with 0 Axes>
<Figure size 640x480 with 0 Axes>
<Figure size 640x480 with 0 Axes>
<Figure size 640x480 with 0 Axes>
<Figure size 640x480 with 0 Axes>
<Figure size 640x480 with 0 Axes>
<Figure size 640x480 with 0 Axes>
<Figure size 640x480 with 0 Axes>
<Figure size 640x480 with 0 Axes>
<Figure size 640x480 with 0 Axes>
<Figure size 640x480 with 0 Axes>
<Figure size 640x480 with 0 Axes>
<Figure size 640x480 with 0 Axes>